

# Lab 4-2 - Diabetic Retinopathy Detection

Department of Computer Science, NCTU

TA 賴佑家

# Important Rules

## Important Date :

- Report Submission Deadline: 4/26 (Tue) **11:59 a.m.(中午11點59分)**
- Demo date: 4/26 (Tue)

## Turn in :

- Experiment Report (.pdf)
- Source code (.py)

**Notice:** zip all files in one file and name it like 「**DLP\_LAB4-2\_your studentID\_name.zip**」, ex: 「**DLP\_LAB4-2\_309553005\_賴佑家.zip**」


# Lab Objective

- In this lab, you will need to analysis diabetic retinopathy (糖尿病所引發視網膜病變) in the following three steps.

**Step 1.** You need to write your own custom DataLoader through PyTorch framework.

**Step 2.** You need to classify diabetic retinopathy grading via the ResNet [1].

**Step 3.** You have to calculate the confusion matrix to evaluate the performance.



# Requirements

- Implement the **ResNet18** 、 **ResNet50** architecture and load parameters from a pretrained model
- Compare and visualize the accuracy trend between the **pretrained model and without pretraining** in same architectures, you need to plot each epoch accuracy (not loss) during training phase and testing phase.
- Implement your own custom **DataLoader**
- Calculate the **confusion matrix** and plotting

# Dataset - Diabetic Retinopathy Detection (kaggle)

- Diabetic retinopathy is the leading cause of blindness in the working-age population of the developed world.
- This dataset provided with a large set of high-resolution retina images taken under a variety of imaging conditions. **Format: .jpeg**



## Class

**0 - No DR**

**1 - Mild**

**2 - Moderate**


**3 - Severe**

**4 - Proliferative DR**



- *Reference : <https://www.kaggle.com/c/diabetic-retinopathy-detection#description>*

# Prepare Data

- **28,099** images for training
- **7025** for testing
- The image resolution is 512x512 and has been preprocessed.
- **Download link:**  
<https://drive.google.com/open?id=1RTmrk7Qu9IBjQYLczaYKOvXaHWBS0o72>
- **Input: [B, 3, 512, 512]**      **Output: [B, 5]**      **Ground truth: [B]**



4752 x 3168



512 x 512

# Dataloader

- Implement your own custom DataLoader
- Below is the skeleton that you have to fill to have a custom dataset, refer to “dataloader.py”

```
class RetinopathyLoader(data.Dataset):  
    def __init__(self, root, mode):  
  
    def __len__(self):  
        """return the size of dataset"""  
  
    def __getitem__(self, index):  
        """something you should implement here"""
```

# Prepare Data

```
def __init__(self, root, mode):  
    """  
    Args:  
        root (string): Root path of the dataset.  
        mode : Indicate procedure status(training or testing)  
  
        self.img_name (string list): String list that store all image names.  
        self.label (int or float list): Numerical list that store all ground truth label values.  
    """  
    self.root = root  
    self.img_name, self.label = getData(mode)  
    self.mode = mode  
    print("> Found %d images..." % (len(self.img_name)))  
  
def __len__(self):  
    """return the size of dataset"""  
    return len(self.img_name)
```



# Prepare Data

- test\_img.csv
- test\_label.csv
- train\_img.csv
- train\_label.csv

3798_left	0
9317_right	0
1991_right	0
2086_left	0
34952_left	0
18072_right	0
9958_left	0
32121_left	0
29612_left	0
21978_left	1
26746_left	0
21469_right	2
40812_right	0
22575_right	2

```
def getData(mode):  
    if mode == 'train':  
        img = pd.read_csv('train_img.csv')  
        label = pd.read_csv('train_label.csv')  
        return np.squeeze(img.values), np.squeeze(label.values)  
    else:  
        img = pd.read_csv('test_img.csv')  
        label = pd.read_csv('test_label.csv')  
        return np.squeeze(img.values), np.squeeze(label.values)
```

**Image Format: .jpeg**



**Please do not sort !!!**

# Prepare Data

```
def __getitem__(self, index):  
    """something you should implement here"""  
  
    """  
    step1. Get the image path from 'self.img_name' and load it.  
           hint : path = root + self.img_name[index] + '.jpeg'  
  
    step2. Get the ground truth label from self.label  
  
    step3. Transform the .jpeg rgb images during the training phase, such as resizing, random flipping,  
           rotation, cropping, normalization etc. But at the beginning, I suggest you follow the hints.  
  
           In the testing phase, if you have a normalization process during the training phase, you only need  
           to normalize the data.  
  
           hints : Convert the pixel value to [0, 1]  
                   Transpose the image shape from [H, W, C] to [C, H, W]  
  
    step4. Return processed image and label  
    """  
  
    return img, label
```



# ResNet

- ResNet (Residual Network) is the Winner of ILSVRC 2015 in image classification, detection, and localization, as well as Winner of MS COCO 2015 detection, and segmentation




# ResNet

- To solve the problem of vanishing/exploding gradients, a skip / shortcut connection is added to add the input  $x$  to the output after few weight layers as below



(a) without skip connections



(b) with skip connections

Figure 1: The loss surfaces of ResNet-56 with/without skip connections. The proposed filter normalization scheme is used to enable comparisons of sharpness/flatness between the two figures.

# ResNet


- ResNet can avoid vanishing gradient problem

$$\begin{array}{ccccccc} x & \longrightarrow & w_1 & \longrightarrow & w_2 & \longrightarrow & w_3 & \longrightarrow & w_4 & \longrightarrow & Loss \\ & & y_1 & & y_2 & & y_3 & & y_4 & & \end{array}$$

$$\begin{aligned} \frac{\partial Loss}{\partial w_1} &= \frac{\partial Loss}{\partial y_4} \frac{\partial y_4}{\partial z_4} \frac{\partial z_4}{\partial y_3} \frac{\partial y_3}{\partial z_3} \frac{\partial z_3}{\partial y_2} \frac{\partial y_2}{\partial z_2} \frac{\partial z_2}{\partial y_1} \frac{\partial y_1}{\partial z_1} \frac{\partial z_1}{\partial w_1} \\ &= \frac{\partial Loss}{\partial y_4} \sigma'(z_4) w_4 \sigma'(z_3) w_3 \sigma'(z_2) w_2 \sigma'(z_1) x_1 \end{aligned}$$

# ResNet

- ResNet can avoid vanishing gradient problem




$$\mathbf{x}_L = \mathbf{x}_l + \sum_{i=l}^{L-1} \mathcal{F}(\mathbf{x}_i, \mathcal{W}_i),$$

$$\frac{\partial \mathcal{E}}{\partial \mathbf{x}_l} = \frac{\partial \mathcal{E}}{\partial \mathbf{x}_L} \frac{\partial \mathbf{x}_L}{\partial \mathbf{x}_l} = \frac{\partial \mathcal{E}}{\partial \mathbf{x}_L} \left( 1 + \frac{\partial}{\partial \mathbf{x}_l} \sum_{i=l}^{L-1} \mathcal{F}(\mathbf{x}_i, \mathcal{W}_i) \right).$$


# ResNet

- ResNe18(Basic block), ResNet50(Bottleneck block)

Basic block



Bottleneck block



# Using Pretrained Model

- Using pretrained model by torchvision module

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    :
  )
  (layer2): Sequential(
    :
  )
  (layer3): Sequential(
    :
  )
  (layer4): Sequential(
    :
  )
  (avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
  (fc): Linear(in_features=512, out_features=1000, bias=True)
)
```


**You need to reinitialize  
the specific layers**



# Result Comparison

- Compare and visualize the accuracy trend **between the pretrained model and without pretraining in same architectures**, you need to plot each epoch accuracy (not loss) during training phase and testing phase.

Result Comparison(ResNet18)



# Confusion Matrix


- A confusion matrix is a table that is often used to describe the performance of a classification model
- `y_true` : ground truth label array
- `y_pred`: prediction array
- `Classes`: label name ['0', '1', '2', '3', '4']

```
def plot_confusion_matrix(y_true, y_pred, classes,  
                           normalize=False,  
                           title=None,  
                           cmap=plt.cm.Blues):
```

*Reference: [https://scikit-learn.org/stable/auto\\_examples/model\\_selection/plot\\_confusion\\_matrix.html#sphx-glr-auto-examples-model-selection-plot-confusion-matrix-py](https://scikit-learn.org/stable/auto_examples/model_selection/plot_confusion_matrix.html#sphx-glr-auto-examples-model-selection-plot-confusion-matrix-py)*

# Confusion Matrix

- Calculate the confusion matrix and plotting



# Hyper Parameters

- Batch size= 4
- Learning rate =  $1e-3$
- Epochs = **10** (resnet18), **5** (resnet50)
- Optimizer: SGD    Momentum = 0.9    Weight\_decay =  $5e-4$
- Loss function: `torch.nn.CrossEntropyLoss()`
- **You can adjust the hyper-parameters according to your own ideas.**
- If you use “`nn.CrossEntropyLoss`”, don’t add softmax after final fc layer because this criterion combines `LogSoftMax` and `NLLLoss` in one single class.

# Report Spec

## Report Spec

1. Introduction (20%)
2. Experiment setups (30%)
  - A. The details of your model (ResNet)
  - B. The details of your Dataloader
  - C. Describing your evaluation through the confusion matrix
3. Experimental results (30%)
  - A. The highest testing accuracy
    - Screenshot
    - Anything you want to present
  - B. Comparison figures
    - Plotting the comparison figures  
**(RseNet18/50, with/without pretraining)**
4. Discussion (20%)
  - A. Anything you want to share

- ----- Criterion of result (40%) -----
- Accuracy  $\geq 82\%$  = 100 pts
- Accuracy 80~82% = 90 pts
- Accuracy 75~80% = 80 pts
- Accuracy  $< 75\%$  = 70 pts
- **Score: 40% experimental results + 60% (report+ demo score)**
- **P.S If the zip file name or the report spec have format error, it will be penalty (-5).**

# Reference

*[1] He, Kaiming, et al. "Deep residual learning for image recognition." Proceedings of the IEEE conference on computer vision and pattern recognition. 2016.*