

NYCU 2022 Spring DLP Lab7 – Let's Play GANs

TA 陳鵬宇


May 17, 2022

Outline

- Lab Objective
- Important Date
- Lab Description
- Scoring Criteria

Lab Objective

- In this lab, you need to implement a **conditional GAN** and generate synthetic images based on multi-labels conditions
- Example of labels:
 - ["cyan cylinder", "red cube"], ["green sphere"], ...



Important Date

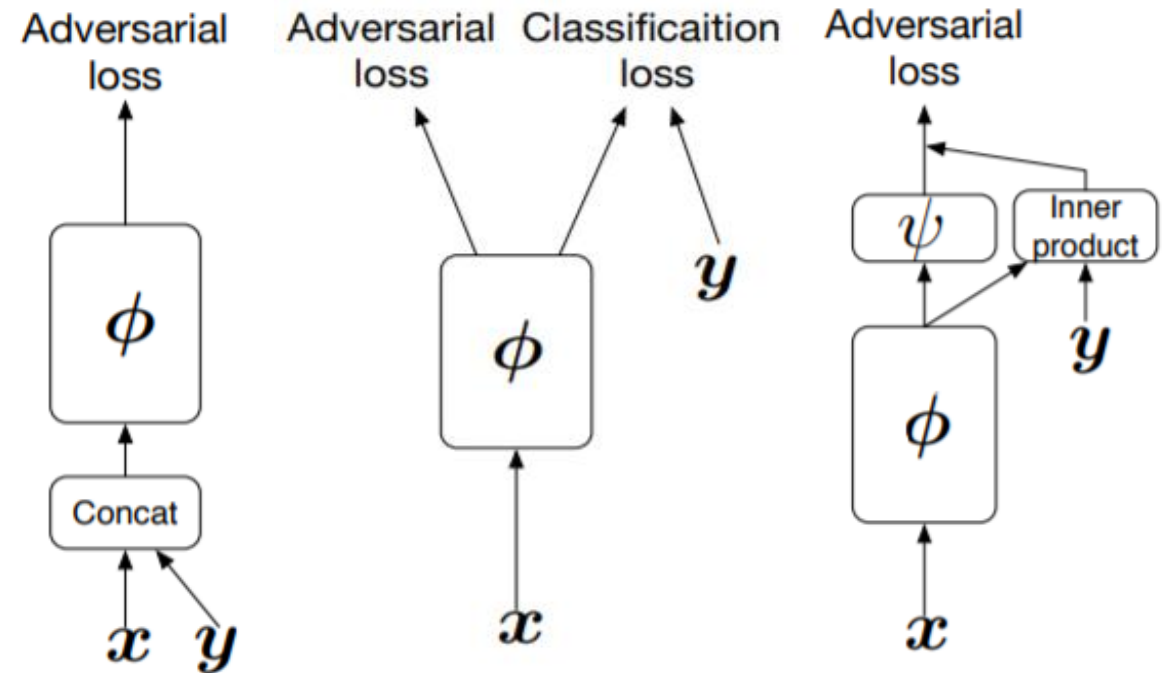
- Experiment Report Submission Deadline: 6/7 (Tue.) 11:59 a.m.
- Demo date: 6/7 (Tue.)
- Zip all files in one file
 - Report (.pdf)
 - Source code
- Name it like "DLP_LAB7_yourstudentID_name.zip"
 - Example: "DLP_LAB7_309551113_陳鵬宇.zip"
- -5% to your score if you do not follow the format

Lab Description

- Implementation details
 - Choose your conditional GAN architecture
 - Design your generator and discriminator
 - Choose your loss function

Lab Description – Choice of cGAN

- Generator
 - Concatenation, multiplication, batch normalization, etc.
- Discriminator
 - Conditional GAN
 - InfoGAN
 - Auxiliary GAN
 - Projection discriminator
- Hybrid version



Lab Description – Design of GAN

- De-convolution layers
- Basic block
- Bottleneck block
- Residual
- Self-attention
- Again, hybrid version
- E.g.
 - DCGAN
 - SA-GAN
 - Progressive GAN

Lab Description – Choice of loss functions





- GAN loss function
 - $L_D^{GAN} = E[\log(D(x))] + E[\log(1 - D(G(z)))]$
 $L_G^{GAN} = E[\log(D(G(z)))]$
- LSGAN
- WGAN
- WGAN-GP
- Combine with your choice of cGAN

Lab Description – Other details

- You can use any GAN architecture your like
- Use the function of a pretrained classifier, **eval**(images, labels), to compute accuracy of your synthetic images.
 - Labels should be one-hot vector. E.g. `[[1,1,0,0,...],[0,1,0,0,...],...]`
 - Images should be all generated images. E.g. (batch size, 3, 64, 64)
- Use **make_grid**(images) and **save_image**(images, path) (from `torchvision.utils import save_image, make_grid`) to save your image (8 images a row, 4 rows)
- The resolution of input for pretrained classifier is 64x64. You can design your own output resolution for generator and resize it.
- Some tips: <https://github.com/soumith/ganhacks>.


Lab Description - Dataset

- Provided files
 - readme.txt, train.json, test.json, object.json, iclevr.zip, evaluator.py, checkpoint.pth
- iclver.zip
 - On the open source google drive
- object.json
 - Dictionary of objects
 - 24 classes




Lab Description – Output examples

Accuracy: 0.667



Accuracy: 0.847



Lab Description – Requirements

- Implement training, testing functions, and dataloader
- Choose your cGAN architecture
- Design your generator and discriminator
- Choose your loss functions
- Output the results based on test.json and new_test.json (will be released before demo)

Scoring Criteria

- Report (40%)
 - Introduction (5%)
 - Implementation details (15%)
 - Describe how you implement your model, including your choice of cGAN, model architectures, and loss functions. (10%)
 - Specify the hyperparameters (learning rate, epochs, etc.) (5%)
 - Results and discussion (20%)
 - Show your results based on the testing data. (5%)
 - Discuss the results of different models architectures. (15%) **For example, what is the effect with or without some specific loss terms, or what kinds of condition design is more effective to help cGAN**

Scoring Criteria

- Demo (60%)
 - Classification accuracy on test.json and new_test.json. (20% + 20%)

• Score ≥ 0.8	----	100%
• $0.8 > \text{score} \geq 0.7$	----	90%
• $0.7 > \text{score} \geq 0.6$	----	80%
• $0.6 > \text{score} \geq 0.5$	----	70 %
• $0.5 > \text{score} \geq 0.4$	----	60 %
• $\text{score} < 0.4$	----	0%
 - Questions (20%)